UZH.ai Hackathon

RepRadar

Listen closely.
Decide confidently.

A farmer’s question.
A team’s next decision.

Syngenta challenge concept · Synthetic-data demonstration



WHY / FICTIONAL ROLE-PLAY

08:12 in Basel. A question arrives from China.

“A headline says some products are already banned. Do I need to change my plans for next season?”

LIN

Farmer in China
Needs practical clarity

MAYA

Corporate Affairs, Basel
Needs a checked response

The missing link: what applies to Lin’s crop, product and location?

Fictional personas and policy scenario. A European proposal does not establish a restriction in China.

Maya has the signals. She still needs the context.

Syngenta works in crop protection and seeds.
Corporate Affairs connects the business with people, public debate and policy.

01	**Three disconnected fragments**	A European proposal, an English ban headline and a Chinese planning question.
02	**One missing local answer**	Which crop, product and market? A translation cannot establish applicability.
03	**A time-sensitive handoff**	Maya must find the right local expert and coordinate a responsible reply.

Lin needs clarity. Maya needs control.

01 **Hear the actual need**

Preserve the request for practical help, even when a post has low reach.

02 **Check before responding**

Separate the shared headline from verified scope and local context.

03 **Keep a human owner**

Maya chooses the next step. Local specialists review technical guidance.

Design goal: help people respond calmly, with clearer information and control.

WHAT / SOLUTION & USER JOURNEY

One question, carried through the whole workflow

01

Listen

Read Lin’s words and the headline.

02

Interpret

Check the proposal and missing local scope.

03

Decide

Use role-specific sources to compare next steps.

04

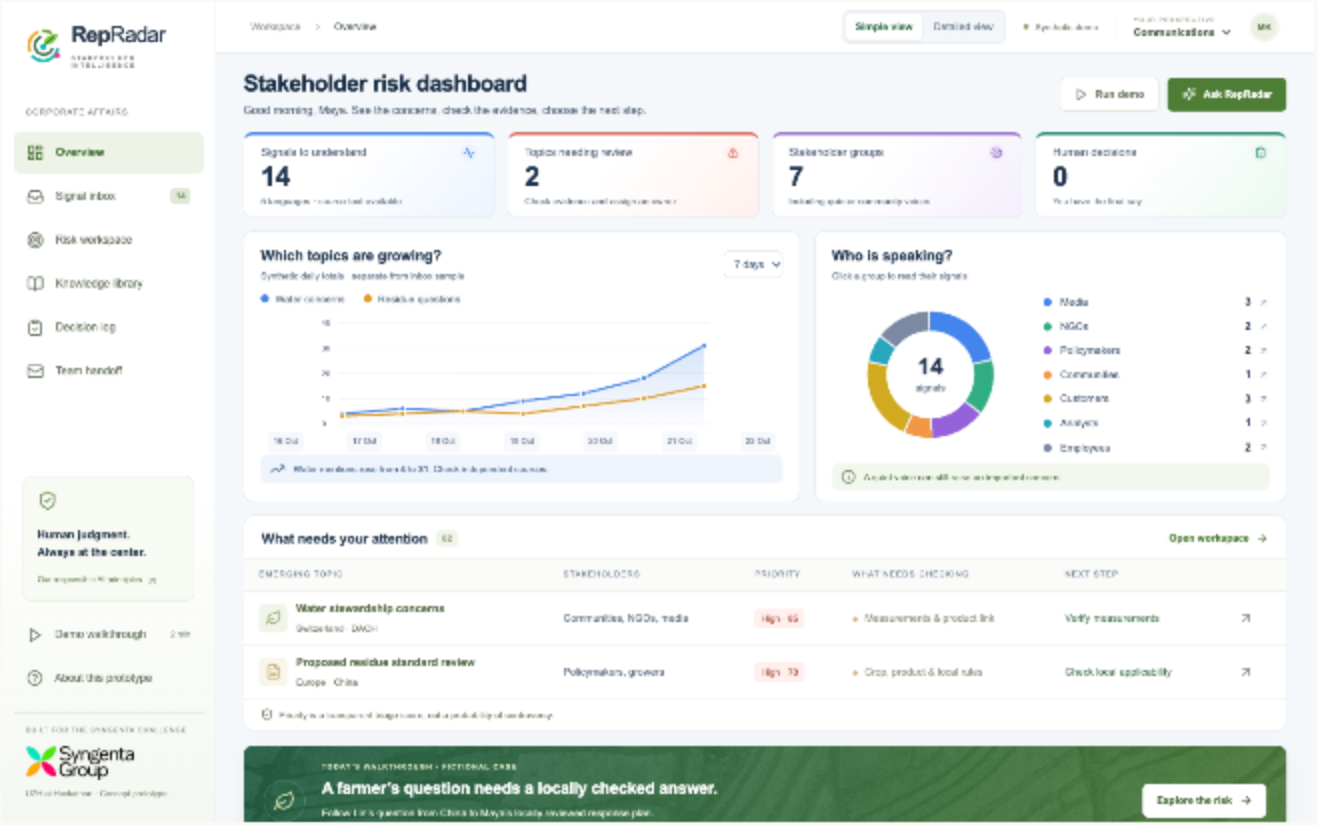
Review

Assign the local owner. Record why and when.

Maya’s decision: verify locally, assign the response owner, and set the next update.

WHAT / TWO-MINUTE WALKTHROUGH

Live demo: help Maya answer Lin



Open live demo

repradar-phi.vercel.app

1 **Read the dashboard**
Open Lin's original question

2 **Check the evidence**
Ask the role-specific assistant

3 **Record a human decision**
Prepare the team handoff

WHAT / AI ARCHITECTURE

Role-specific RAG and a human review boundary

Stage	Working prototype	Pilot extension
Signals	Multilingual synthetic records	Licensed feeds, redaction, deduplication
Interpretation	Authored agent perspectives	Media, regulatory and local-language agents
Knowledge	Role-filtered retrieval (RAG context)	Permission-filtered hybrid / vector retrieval
Recommendations	Ranking; optional local Qwen3	Evaluated generation with source tracing
Decision	Human review + Outlook / CSV	SSO, mailbox and ERP connectors

Context: reviewed policies, agricultural knowledge, historical cases and expert interviews.

Maya can disagree. Lin’s context stays visible.

01 Agency

Maya can edit, reject or defer. Every decision needs a reason and owner.

02 Cultural meaning

Read Lin’s actual question. A local colleague checks meaning and technical scope.

03 Privacy and fairness

Use licensed or consented sources. Avoid profiling people. Include quieter voices.

Pilot safeguards: SSO, document permissions, retention rules and an approved model environment.

LEARNING / CONTROLLED IMPROVEMENT

The story continues after the decision

01

Decision

Maya assigns a China owner and review time.

02

Outcome

The local team checks whether the reply helped.

03

Curator review

A curator checks the observed lesson.

04

Knowledge update

The next team retrieves the approved lesson.

Outcome follow-up is a proposed workflow. The demo records decisions and reviewed lessons locally.

FEASIBILITY / ADOPTION & COST

A six-week pilot with ten Corporate Affairs colleagues

1–2 Co-design the workflow

Two regions, two topics. Interview employees and establish the current baseline.

3–4 Run in shadow mode

Compare alerts with expert review. No autonomous publication.

5–6 Evaluate and decide

Review usability, missed concerns and citation quality before wider adoption.

Illustrative operating budget: about CHF 3.4k/month, plus integration and setup.

Planning assumptions, not a vendor quote. Human review and data access dominate this example.

One complete story demonstrates the five criteria

Criterion	Evidence in the proposal
Problem–solution fit	Dashboard → scoped evidence → response plan → Outlook handoff
Feasibility	Local Qwen3, editable draft / CSV, six-week regional pilot
Originality	Local context + role-specific RAG + preferences + reviewed lessons
Human-centered & responsible	Simple view, quieter voices, visible uncertainty, edit / reject / review
Pitch & Q&A	The same farmer, employee and question from opening to closing

Value to test: less briefing effort and better-supported decisions, with human control intact.

RepRadar

Back to Lin.
A clear next step.

A plan to verify. A local owner.
A scheduled update, with human review.

Our ask

One sponsor. Ten colleagues.
Two regions. Six weeks.

Listen closely. Decide confidently.



APPENDIX / COST MODEL

Pilot economics: assumptions you can change

Monthly item	Assumption	CHF / month
Model inference	2,200 requests × 8k input / 1k output tokens	79.20
Application hosting	Illustrative allowance	300.00
Data access	Illustrative licensed-source allowance	1,000.00
Human review	20 hours × CHF 100/hour	2,000.00
Operating total	Excludes setup, legal work and local-model hardware	3,379.20

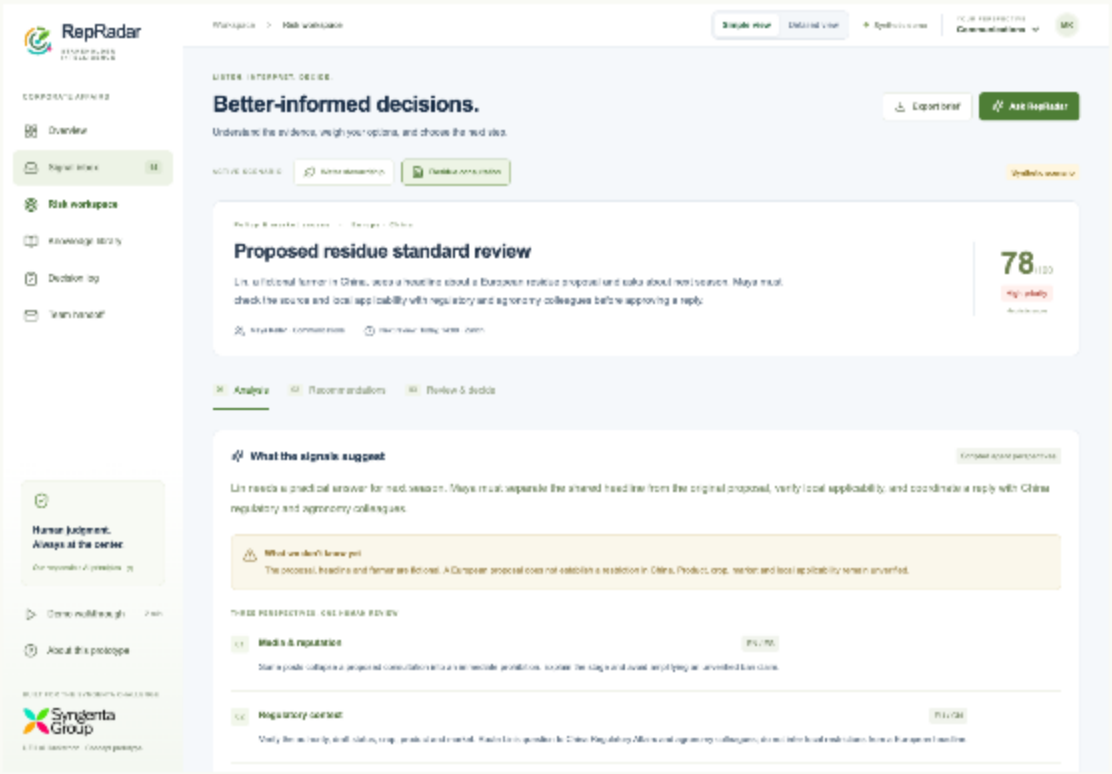
Token rate assumptions: CHF 3/M input and CHF 12/M output. These are hypothetical rates.

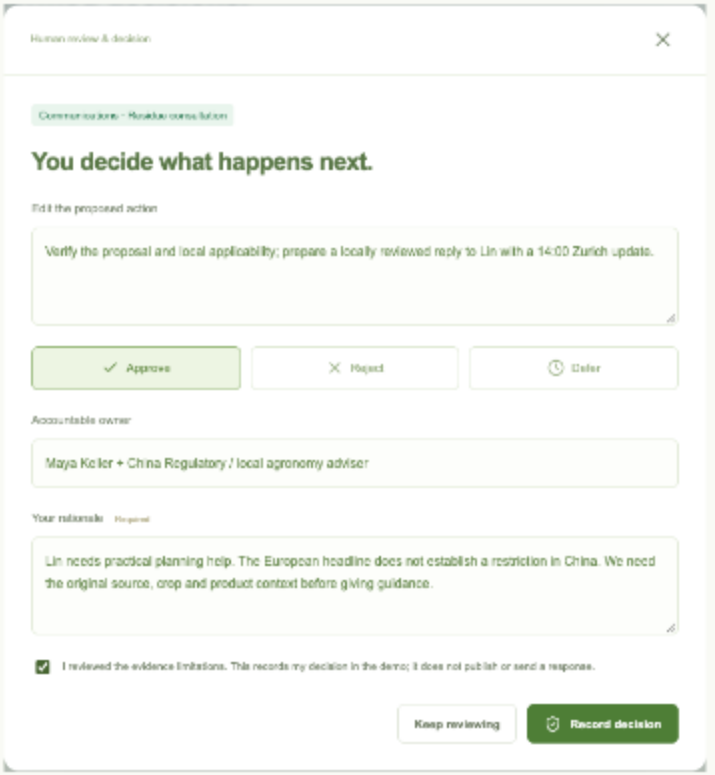
RepRadar

Syngenta Group

APPENDIX / OFFLINE FALLBACK

Demo backup: Maya’s response plan for Lin





Questions we expect

Q Does it predict controversy?	It prioritizes early signals. Validated probability forecasts need labelled history.
Q Why not use a generic chatbot?	Role-filtered sources, original language and a recorded decision give the work structure.
Q Is it ready for internal data?	The demo uses synthetic data. Identity, permissions and approved hosting come first.

Next validation: local users, source coverage, false alarms, missed concerns and employee trust.

Demo coverage: understand the concern

What the employee can do	Why it matters	Evaluation criterion
Scan the dashboard	Overview plus sentiment, languages, risk drivers and source / review charts	Fit · Human-centered
Filter and inspect signals	Search, group / language filters, original text, translation and review notes	Fit · Human-centered
Inspect risk and uncertainty	Visible factors, editable weights, evidence gaps and a simulated new signal	Originality · Human-centered
Inspect role-specific knowledge	Document IDs, versions, owners and role / scenario filtering	Fit · Feasibility
Ask RepRadar in EN / ZH	Source excerpts online; local Qwen3 option, chat context and source links	Originality · Feasibility

Sentiment, translations and agent perspectives are authored demo examples; live ingestion is a pilot step.

Demo coverage: decide, coordinate and learn

What the employee can do	Why it matters	Evaluation criterion
Compare three next steps	Change evidence, trust and speed preferences; inspect the ranking	Originality · Human-centered
Edit, approve, reject or defer	Own the action, record a rationale and acknowledge evidence limits	Fit · Human-centered
Prepare a team handoff	Editable Outlook compose link, copied brief or reviewed ERP CSV	Fit · Feasibility
Keep an accountable record	Decision log, export and browser-local persistence	Feasibility · Human-centered
Review outcomes and lessons	Capture feedback; curator approval before role-specific reuse	Originality · Human-centered
Rehearse and recover	Guided walkthrough, reset, source fallback and deck screenshots	Feasibility · Pitch & Q&A

Outlook: human sends. ERP: CSV export. Learning: curated retrieval, not model retraining.